

Do companies committing accounting fraud write their annual reports differently?

A pre-registered novelty study on SEC 10-K filings.

Trained: five small autoencoders in PyTorch, one per held-out fraud cohort.

Reused: a pretrained MiniLM transformer, frozen, as a sentence-to-vector encoder.

Tested: six historical accounting frauds; methodology locked in git before any number was looked at.





Does an unsupervised autoencoder trained on industry-and-year-matched clean 10-K MD&A text assign elevated reconstruction error to known historical fraud filings, under strictly out-of-sample evaluation?

The one-line answer this report defends:

At $N = 6$ frauds, with strict pre-registration: no useful autoencoder signal across four of five evaluable cohorts; Enron is not evaluable under the frozen rule; the one cohort the autoencoder did rank highly (Lehman) is $n = 1$, and a 1990s TF-IDF baseline beats the autoencoder on it.

The contract is the contribution.

git-tagged validation-spec-frozen, committed before any validation script ran.

Out-of-sample, always

Leave-one-cohort-out (LOCO): for each held-out fraud, the autoencoder is trained on clean filings from other cohorts only. The fraud and its peers are never in the training set.

Time-controlled training

Within the LOCO training set, only filings dated on or before that fraud's filing date are eligible. The autoencoder's training corpus cannot include post-fraud disclosure language. (MiniLM pretraining is a separate caveat.)

Single pass, no tuning

Architecture, hyperparameters, sentence cap, seeds: all committed and tagged before validation ran.
scripts/o6_validate.py runs once. No model tuning after validation.

Whatever the test produces, the contract makes the result interpretable.

git tag: **validation-spec-frozen**
at 98d2585 · 2026-05-01 13:42:42 EDT

Only Lehman ranks better than random expectation; even there, a 1990s baseline beats the autoencoder.

Cohort	Rank	Note
Lehman (LEH, FY 2007) [△]	3 / 13	<i>n = 1, baseline beats autoencoder</i>
HealthSouth (HRC)	7 / 12	at or below random
Valeant (VRX)	8 / 13	at or below random
Tyco (TYC)	9 / 13	at or below random
WorldCom (WCOM)	12 / 13	at or below random
Enron (ENE, FY 2000)	—	<i>not evaluable: filed earliest, no eligible training data</i>

Aggregate hit-rates

hit@1: **0.00** vs random 0.08


hit@3: **0.20** vs random 0.23

hit@5: **0.20** vs random 0.39

[△] *Lehman is the only rank better than random expectation, and a 1990s baseline beats the autoencoder there too (slide 5). The aggregate is entirely Lehman.*

A 1990s baseline matched or beat the autoencoder on every cohort.

Post-hoc: tests whether the autoencoder earned its complexity, does not rescue the primary result.



How to read the chart: shorter bar = lower rank number = better detection (rank 1 is most anomalous). Lehman: autoencoder rank 3/13, TF-IDF rank 1/13. Aggregate hit@5 is 0.40 (TF-IDF) vs 0.20 (autoencoder). The pre-registered model does not beat the trivial baseline.

Q Did you check a trivial baseline?

Yes. TF-IDF + truncated-SVD, same bottleneck dimension, same training corpora. It matches or beats the autoencoder on every cohort. The honest claim is that the autoencoder adds nothing here. Slide 5.

Q The baseline is post-hoc. Isn't that HARKing (hypothesizing after results)?

The pre-registered hypothesis was about the autoencoder, and the autoencoder result is reported unchanged. The TF-IDF baseline was added after seeing the null and is reported as post-hoc. The primary null does not depend on it.

Q Could MiniLM have memorized 'Lehman' / 'Repo 105' / etc.?

Tested post-hoc. Masked 93 such mentions with [ENTITY] and re-scored. Lehman's rank held at 3/13 (delta = 0). Exact-token name leakage did not explain the rank; topical leakage remains open.

CONCLUSION

What survived: pre-registration.
What failed: the autoencoder beyond the TF-IDF baseline.

A result that loses to a baseline is not evidence for the neural model. The contract held; the model did not.

What I'd do differently

Confirm every held-out fraud is evaluable under the pre-registered rule BEFORE freezing the spec. Enron taught that the hard way.

Next experiment

Same autoencoder, within-firm temporal delta. Does Lehman's MD&A look anomalous against Lehman's own prior-year MD&A?

Pipeline demo (not a fraud scanner): canary-psi.vercel.app/scan

Code: github.com/ArseniiChan/canary · Tag: [validation-spec-frozen](#)

Thanks.

Scan to try the pipeline

